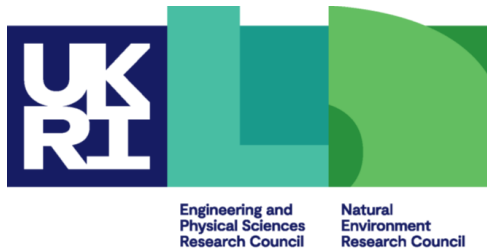


# Advanced Message-Passing Programming

---

## Overview of Parallel IO



# Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

<http://creativecommons.org/licenses/by-nc-sa/4.0/>

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.

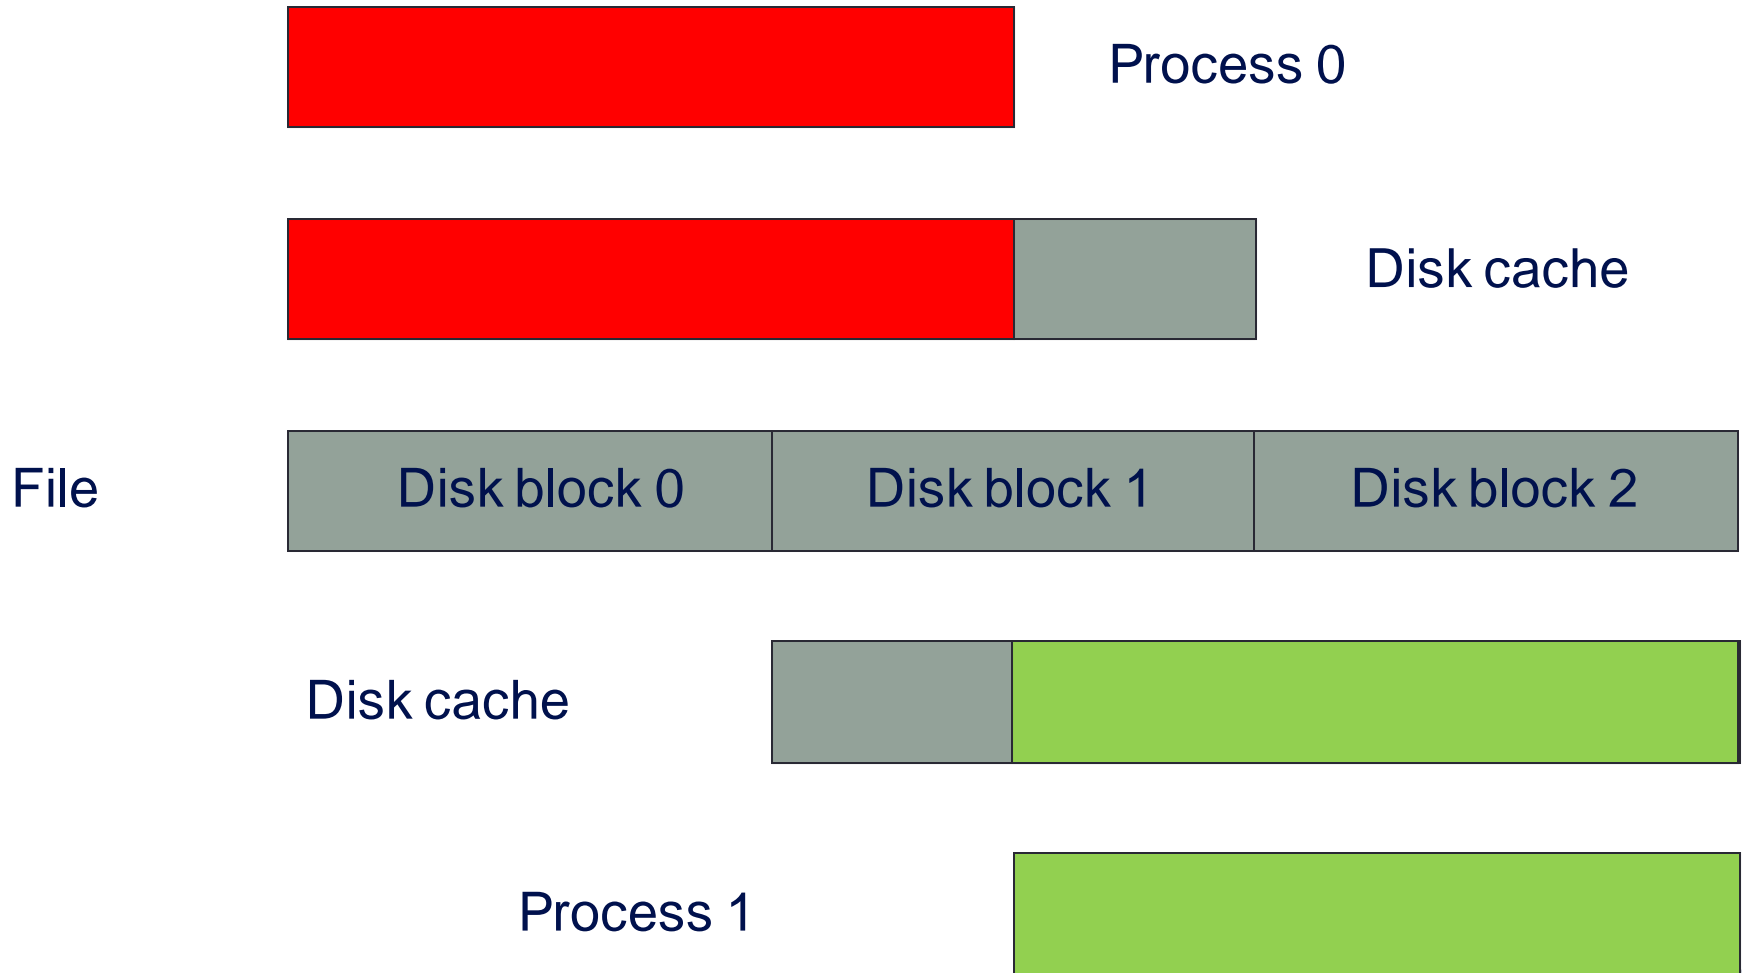
# Why is IO hard?

- Breaks out of the nice process/memory model
  - data in memory has to physically appear on an external device
- Files are very restrictive
  - linear access probably implies remapping of program data
  - just a string of bytes with no memory of their meaning
- Many, many system-specific options to IO calls
- Different formats
  - text, binary, big/little endian, Fortran unformatted, HDF5, NetCDF, ...
- Disk systems are very complicated
  - RAID disks, many layers of caching on disk, in memory, ...
- IO is the HPC equivalent of printing!

# Why is Parallel IO Harder?

- Cannot have multiple processes writing to a single file
  - Unix generally cannot cope with this
  - data cached in units of disk blocks (eg 4K) and is *not coherent*
  - not even sufficient to have processes writing to distinct parts of file
- Even reading can be difficult
  - 1024 processes opening a file can overload the filesystem (fs)
- Data is distributed across different processes
  - processes do not in general own contiguous chunks of the file
  - cannot easily do linear writes
  - local data may have halos to be stripped off


# Simultaneous Access to Files



# Parallel File Systems

- Parallel computer
  - constructed of many standard processors, each not particularly fast
  - performance comes from using many processors at once
  - requires manual distribution of data and calculation across processors
- Parallel file systems
  - constructed from many standard disks, each not particularly fast
  - performance comes from reading / writing to many disks at once
  - requires many *clients* to read / write to different disks
    - each *node* appears as a separate IO client
  - data from a single file can be *striped* across many disks
- Must appear as a single file system to user
  - typically have a single *MetaData* Server (MDS)

# Parallel File Systems: Lustre



# ARCHER's (not ARCHER2) Cray Sonexion Storage

## MMU: *Metadata Management Unit*



### Lustre MetaData Server

- Contains server hardware and storage

## SSU: *Scalable Storage Unit*



2 x OSSs and 8 x OSTs (Object Storage Targets)

- Contains Storage controller, Lustre server, disk controller and RAID engine
- Each unit is 2 OSSs each with 4 OSTs of 10 (8+2) disks in a RAID6 array



Multiple SSUs are combined to form storage racks

3




# Independent I/O

- Also called File-per-process
- By default, each file stored on a single OST
  - simplest to think of each Object Storage Target as a disk
- The MDS will assign different files to different disks
  - in principle achieves maximum bandwidth
- But ...
  - a nightmare to cope with in practice (thousands of files)
  - opening and closing files is serialised via a single MDS
    - this can be a performance bottleneck
- Mitigations
  - multiple filesystems with fewer OSTs, each with their own MDS
  - write a file-per-node rather than file-per-process

# Lustre data striping

Parallel performance comes from striping single files over multiple OSTs



Single logical user  
file e.g.  
/work/q01/q01/user  
/bigfile.dat

OS/file-system  
divides the file into  
stripes *if requested  
by the user*

Stripes read/written to/from  
their assigned OST

# Parallel File Systems


- Allow multiple IO processes to access same file
  - increases bandwidth
- Typically optimised for bandwidth
  - not for latency
  - e.g. reading/writing small amounts of data is very inefficient
- Very difficult for general user to configure and use
  - need some kind of higher-level abstraction
  - allow user to focus on data layout across user processes
  - don't want to worry about how file is split across IO servers

# 4x4 array on 2x2 Process Grid

Parallel Data

2	4	2	4
1	3	1	3
2	4	2	4
1	3	1	3

File



1	2	1	2	3	4	3	4	1	2	1	2	3	4	3	4
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

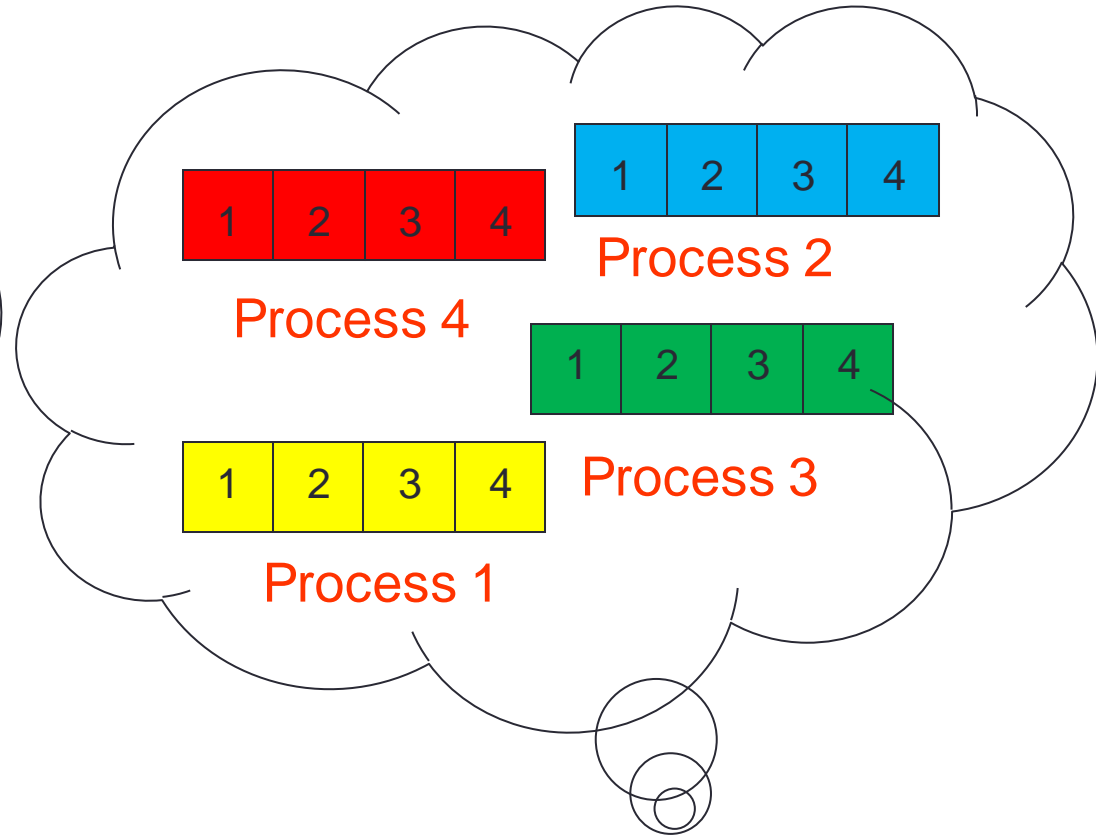
# Message Passing: Naive Solutions

- Controller IO
  - send all data to/from single controller process and write/read a single file
  - quickly run out of memory on the controller
    - or have to write in many small chunks
  - does not benefit from a parallel fs that supports multiple write streams
- Separate files
  - each process writes to a local fs and user copies back to home
  - or each process opens a unique file (dataXXX.dat) on shared fs
- Major problem with separate files is reassembling data
  - file contents dependent on number of CPUs and decomposition
  - pre / post-processing steps needed to change number of processes
  - but at least this approach means that reads and writes are in parallel
- But may overload filesystem for many processes
  - e.g. MDS cannot keep up with requests

# Programmer View vs Machine View



4	8	12	16
3	7	11	15
2	6	10	14
1	5	9	13



1	2	3	4
---	---	---	---

Process 4

1	2	3	4
---	---	---	---


Process 2

1	2	3	4
---	---	---	---

Process 3

1	2	3	4
---	---	---	---

Process 1



# MPI-IO Approach

- MPI-IO is part of the MPI standard
  - <http://www.mpi-forum.org/docs/docs.html>
- Each process needs to describe what subsection of the global array it holds
  - it is entirely up to the programmer to ensure that these do not overlap for write operations!
- Programmer needs to be able to pass system-specific information
  - pass an `info` object to all calls

# Data Sections

4	8	12	16
3	7	11	15
2	6	10	14
1	5	9	13

on process 3

4	8	12	16
3	7	11	15
2	6	10	14
1	5	9	13

- Describe 2x2 subsection of 4x4 array
- Using standard MPI derived datatypes
- A number of different ways to do this
  - we will cover three methods in the course



# Other Parallel IO Libraries

- MPI-IO is usually the lowest level
  - you may never call it directly
- Higher level libraries exist
  - HDF5
  - NetCDF
  - ADIOS
- Approach is the same
  - some way of describing what portion(s) of file each process owns
  - these libraries usually use MPI-IO to do their reads and writes

# Summary

- Parallel IO is difficult
  - in theory and in practice
- MPI-IO provides a higher-level abstraction
  - user describes global data layout using derived datatypes
  - MPI-IO hides all the system specific filesystem details ...
  - ... but (hopefully) takes advantage of them for performance
- More flexible formats like NetCDF and HDF5 exist
  - they gain performance by layering on top of MPI-IO
- User requires a good understanding of derived datatypes
  - see next lecture